



Benchmarking Redis vs Memcached

Vonteri Harshith Reddy - 2019CS50450

Memcached

- Open source, In-memory data store
- Only Key value pairs
- Client-server architecture
- Multi-threaded
- Allows only bit strings as values
- No support for Persistence



REmote DIctionary Server - Redis

- Open source, In-memory data store
- Primarily key value store but can act as multi-model(document-store)
- Client-server architecture
- Single threaded
- Complex Data Structures
- On-disk Persistence



Why Redis and Memcached?

- Redis and Memcached are in-memory key value stores
- In DB ranking of Key value stores Redis stands 1st and Memcached stands 4th
- Redis stands 1st and Memcached 2nd when only in-premise data stores are considered as Azure and DynamoDB can only be deployed on cloud
- Redis and Memcached are top competitors in this domain
- Why is this domain important? Gaming leaderboards, carts, user sessions etc.,

| RANK(NOV2023) | DBMS | SCORE(NOV2023) |
|---------------|-----------------|----------------|
| 1 | Redis | 160.02 |
| 2 | Microsoft Azure | 83.24 |
| 3 | Amazon DynamoDB | 34.11 |
| 4 | Memcached | 19.8 |



Experimental Setup


- Benchmarking tool: Mem-tier
 - Supports both testing and performance measure
 - Opensource, developed by Redis
 - Compatible with Memcached and Redis
 - Multithreaded
 - Customizable parameters
 - Statistical Analysis
- Redis version - 7.0.12
- Memcached version – 1.6.22
- 8 Core CPU, 8GB RAM

Experiments

- Benchmarked Redis, Memcached for:
 - Latency
 - Throughput
 - Scalability
 - Fault tolerance


Latency Benchmark

- Mem-tier Benchmark
 - Single thread, 1 client per thread
 - Data size: 32 bytes
- Memcached exhibits slightly better performance than Redis on Set and Get operations
- Trends show that Get operation is slightly faster than Set operation
- Redis-Set operation is 5.5% slower than Memcached-Set, Redis-Get operation is 6% slower than Memcached-Get




Throughput Benchmark

- Memtier Benchmark
 - 4 threads, 50 clients per thread
 - Each client sends 1000000 requests
 - Data size: 10 bytes
- Memcached clearly has higher throughput than Redis, as it uses multi-threading for request serving




Throughput Benchmark – Dictionary Expansion

- A threshold is set in Redis and Memcached for controlling when dictionary expansion is necessary
- Once the threshold is exceeded, a double-size hash bucket is created, and data gradually moves from existing bucket to the new one
- 2 dents are observed in the plot for Redis, which implies that the server experienced dictionary expansion twice
- Throughput degradation in Redis due to dictionary expansion is ~38%



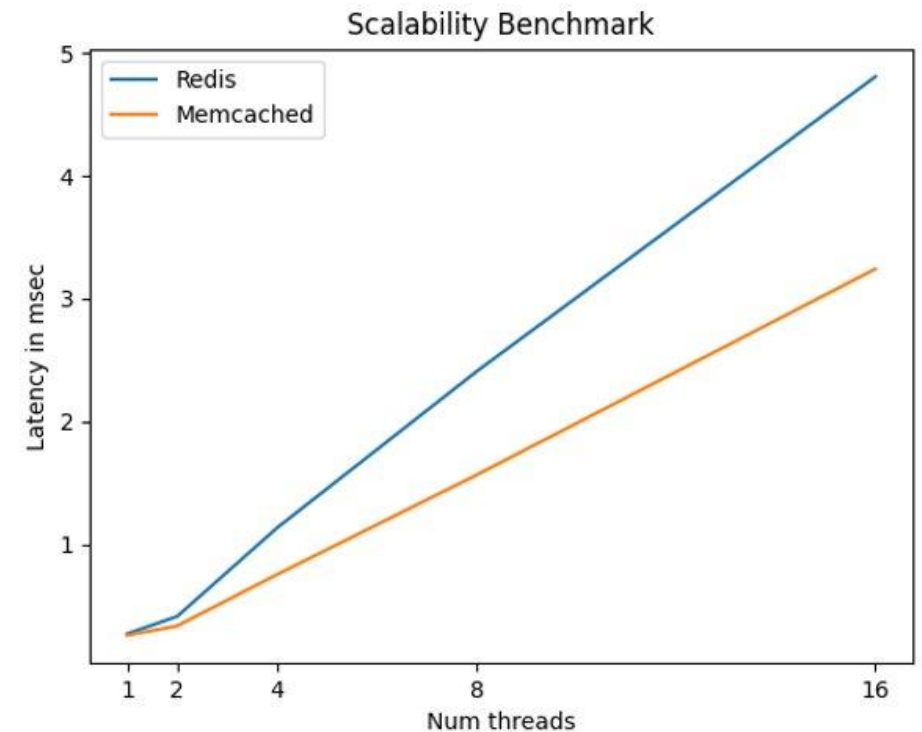
Throughput Benchmark – Dictionary Expansion

- Memcached didn't experience much dent in throughput even with dictionary expansion
- Memcached uses multithreading, dedicates one exclusive expanding thread, always ready to perform dictionary expansion
- In contrast Redis only has single thread which must deal with request serving and dictionary expansion
- The increase in throughput of Memcached may be attributed to a reduction in the number of active clients over time



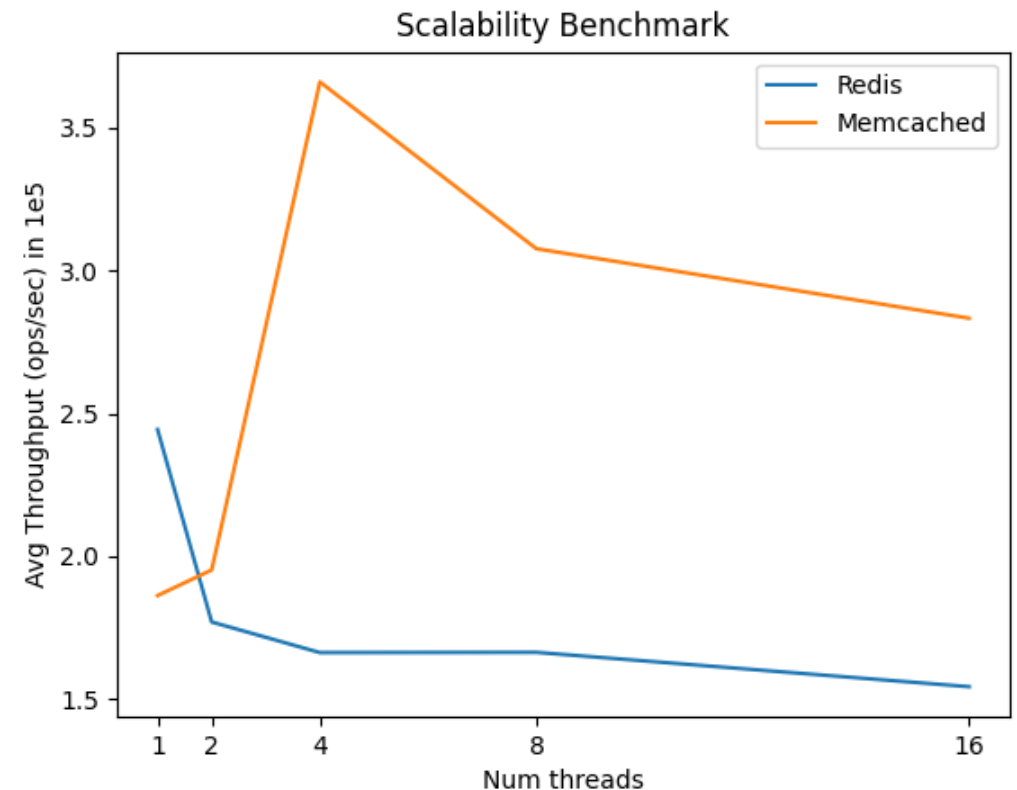
Scalability Benchmark - Latency

- Mementier Benchmark
 - Each client sends 10000 requests
 - Data size: 32 bytes
 - 50 clients per thread
 - Set:Get ratio 1:10
- Memcached clearly exhibits better scalability in Avg. latency with increase in no of threads, thanks to its use of multiple threads to serve requests
- Avg. Latency of Redis is 48.3% more than that of Memcached



Scalability Benchmark - Throughput

- Throughput of Memcached is significantly higher than Redis even with more no of threads
- Redis is single-threaded, hence there is a constant drop in throughput with increase in no of threads
- Because of multi-threaded nature of Memcached, there is an increase in throughput up to 4 threads, a decrease thereafter





Fault Tolerance

- Memcached is not fault-tolerant as it doesn't support on-disk persistence
- Redis is fault tolerant and has two persistence models:
 1. RDB (Redis Database) Persistence: Performs point-in-time snapshots of your dataset at specified intervals
 2. AOF (Append Only File) Persistence: AOF persistence logs every write operation received by the server. These operations can then be replayed again at server startup, reconstructing the original dataset
 - 3 modes for syncing: Always (Immediate), Every Sec, No (OS flush the output buffer when it wants)
- We can disable persistence completely. This is sometimes used when caching.

Overhead of Persistence Models

- We compare the performance of 3 models of Redis, namely RDB, AOF-Always, No Persistence to understand the overhead of Persistence
- 4 threads, 50 clients per thread, 100000 requests per clients
- Negligible overhead of RDB is that Redis forks a child process by using copy-on-write for the background save purpose, thus loading data and snapshotting can run concurrently
- Significant overhead is observed in AOF-Always (Immediate) sync.

| MODEL | AVG. LATENCY (MSEC) | AVG. THROUGHPUT (OPS/SEC) |
|-------------------|---------------------------|---------------------------------|
| AOF-Always | 7.14429 | 28021.47 |
| RDB | 1.13638 | 176008.06 |
| No persistence | 1.12054 | 176351.48 |

Memory Usage

- 1 client, 1 thread
- Only Set operations
- Memory Usage (in MB) of Memcached is better than Redis with varying sizes of database as memcached uses less metadata to store key, values

| Num Keys | 10000 | 100000 | 1000000 |
|-----------|-------|--------|---------|
| Redis | 0.78 | 7.46 | 72.4 |
| Memcached | 0.67 | 6.87 | 49.6 |

Summary – Redis Vs Memcached

- **Latency:** Both are fast with slight edge for Memcached
- **Throughput:** Memcached has higher throughput due to its use of multiple threads
- **Scalability:** Memcached Scales better than Redis (single instance)
- **Persistence, Replication, Complex Data structures:** Only supported in Redis
- **Memory Usage:** Memcached is better

Applications



Redis:

Gaming leaderboards
Ecommerce cart
user info on social media platforms
Search engines for full text search
Chat applications
Geospatial applications



Memcached:

Transient data caching
Stateless Caching for Web Applications

Future Work

- Perform extensive experiments like analyzing latency in terms of 99percentile
- Benchmark Availability
- Benchmark Consistency





Thank You

